

# Deep Learning: Lecture 1

Keith L. Downing

The Norwegian University of Science and Technology (NTNU)  
Trondheim, Norway  
keithd@idi.ntnu.no

January 12, 2020

- 1 Gradient Descent Learning
- 2 Matrix Representations of Neural Networks
- 3 Activation Functions
- 4 Loss Functions
- 5 Maximum Likelihood Estimators

# Gradient-Based Optimization (GBO)

- $L$  = **objective** function.
- In GBO, use partial derivs of  $L$  (w.r.t. system params,  $W$ ) to determine intelligent modifications of  $W$  so as to maximize or minimize  $L$ .

## Simple Example

- $\Psi$  = set of cases; each case is a pair  $(a,b)$
- $W$  = parameters of the system.  $W = (x,y)$
- $L(a,b) = ax + by$  = the objective function
- Goal = Maximize  $\sum_{k \in \Psi} L(a_k, b_k)$
- **Problem:** You will never see complete contents of  $\Psi$ .
- **Solution:** Best alternative = Tune  $W$  to maximize  $L$  for the **samples** you do see. For each case  $(a_j, b_j)$ , nudge  $x$  and  $y$  in directions that increase  $L(a,b)$  for that case. Size of nudge = learning rate =  $\eta$ .
- Mathematically, use gradients:
  - $\Delta x = \eta \frac{\partial L(a,b)}{\partial x} \big|_{(a_j, b_j)} = \eta a_j$
  - $\Delta y = \eta \frac{\partial L(a,b)}{\partial y} \big|_{(a_j, b_j)} = \eta b_j$

# Gradient Descent Learning (GDL)

- $L$  = **loss** or **cost** function.
- In GDL, use partial derivs of  $L$  (w.r.t. system params,  $W$ ) to determine intelligent modifications of  $W$  so as to minimize  $L$ .

## Simple Example

- $\Psi$  = set of cases,  $c_i = (f_{i,1}, f_{i,2}, \dots, f_{i,m}; t_i)$  = features + target value.
- $W$  = parameters (e.g. weights) of the system = a vector.
- $P(f_i) = f_i \bullet W$  = the prediction function
- Goal = Minimize  $L(\Psi) = \text{loss function} = \sum_{k \in \Psi} |P(f_k) - t_k|$
- **Problem:** You will never see complete contents of  $\Psi$ , but you want to make good predictions  $\forall c \in \Psi$
- **Solution:** Best alternative = Tune  $W$  to minimize  $L$  for the **samples** you do see. For each case  $c_k = (f_k, t_k)$ , nudge each  $w_i$  in direction that decreases  $L(c_k)$ . Size of nudge = learning rate =  $\eta$ .
- Mathematically, use gradients:
  - $\Delta w_i = -\eta \frac{\partial L(c)}{\partial w_i} |_{(f_k, t_k)} = \eta \text{ ?????? (see next slide)}$

# Differentiating the Loss Function

Mean Squared Error (MSE) = a common loss function:

$$L(\Psi) = \frac{1}{\|\Psi\|} \sum_{k \in \Psi} (P(f_k) - t_k)^2$$

By the **Chain Rule** of calculus:

$$\frac{\partial L(\Psi)}{\partial w_i} = \frac{1}{\|\Psi\|} \sum_{k \in \Psi} 2(P(f_k) - t_k) \frac{\partial (P(f_k) - t_k)}{\partial w_i} = \frac{2}{\|\Psi\|} \sum_{k \in \Psi} (P(f_k) - t_k) f_{k,i}$$

$$\text{since } \frac{\partial P(f_k)}{\partial w_i} = f_{k,i} \text{ and } \frac{\partial t_k}{\partial w_i} = 0.$$

When we only have a **sample**  $S$  of all the cases in  $\Psi$ , we estimate the partial derivative (a.k.a. gradient) as:


$$\frac{\partial L(S)}{\partial w_i} = \frac{2}{\|S\|} \sum_{k \in S} (P(f_k) - t_k) f_{k,i}$$

And we update each  $w_i \in W$  to try to reduce this loss:

$$\Delta w_i = -\eta \frac{2}{\|S\|} \sum_{k \in S} (P(f_k) - t_k) f_{k,i}$$

The 2 in  $\frac{2}{\|S\|}$  is usually ignored and implicitly incorporated into the chosen value of  $\eta$


# Complexifying Neural Networks



## How do these networks get more complex?

- Nonlinear activation functions (applied to sum of weighted inputs)
- Multiple output neurons
- Multiple layers of neurons separating inputs from outputs.
- Various types of layers
- More elaborate objective (loss) functions.
- .... and much more ....

# Amazingly Broad Applicability of Gradient Descent



- Despite increasingly complex NNs, the essence of gradient descent learning remains the same: for EVERY weight  $w_i$  in the network:

$$\Delta w_i = -\eta \frac{\partial L(c)}{\partial w_i}$$

- Calculating these gradients just gets more complicated.
- ML packages such as Tensorflow, PyTorch and Theano automate this !!

# Two Basic Types of Gradient Descent Learning

## ● Batch Gradient Descent (GD)

- Run ALL available training cases ( $S$ ) (a *batch*) through NN.
- Calculate gradients **for each case** in  $S$ .
- Update weights, but only after entire batch is run and all gradients are combined (e.g. averaged).
- Repeat
  - \* Each processing round (over an entire batch) = an **epoch**


## ● Stochastic Gradient Descent (SGD)

- Run a **subset**  $C$  (often randomly selected) of  $S$  (a **minibatch**) through NN.
- Calculate gradients for each  $c \in C$
- Update weights based on combined gradients.
- Repeat
  - \* Potentially unstable for small  $\|C\|$ .
- \*\* For SGD, an epoch =  $\frac{\|S\|}{\|C\|}$  minibatches.

For GD (SGD), each case in a batch (minibatch) experiences the same weights when run through the net.



# Matrices for Neural Networks



$$\begin{bmatrix} x_1 & x_2 & x_3 \end{bmatrix} \times \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \\ w_{31} & w_{32} \end{bmatrix} =$$

Output Values

$[1 \times 3] \times [3 \times 2] \Rightarrow [1 \times 2]$

Incoming weights      Incoming weights

$$\left[ \begin{array}{c} x_1 w_{11} + x_2 w_{21} + x_3 w_{31} \\ x_1 w_{12} + x_2 w_{22} + x_3 w_{32} \end{array} \right]$$

Sum weighted inputs to 1      Sum weighted inputs to 2

# Standard Representations

$$X = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$

\* Standard notation: activation vectors are COLUMN vectors:

3 x 1 Column Vector

$$\begin{matrix} X^T \\ \begin{bmatrix} x_1 & x_2 & x_3 \end{bmatrix} \\ 1 \times 3 \text{ Row Vector} \end{matrix} \begin{matrix} W \\ \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \\ w_{31} & w_{32} \end{bmatrix} \\ 3 \times 2 \text{ Matrix} \end{matrix} = \begin{matrix} X^T W \\ 1 \times 2 \text{ Row Vector} \end{matrix}$$

$$X^T W = (W^T X)^T$$

# Transposed Weight Matrices

Transpose all weight matrices (just once) and then use  $W^T X$  to produce new **column** vectors. Then activation vectors need not be transposed back and forth between column and row vectors.

$$\begin{matrix} & W^T & & X_0 \\ \left[ \begin{array}{ccc} w_{11} & w_{21} & w_{31} \\ w_{12} & w_{22} & w_{32} \end{array} \right] & \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} & = & W^T X_0 \end{matrix}$$

$2 \times 3$  Matrix                       $3 \times 1$  Column Vector                       $2 \times 1$  Column Vector

$2 \times 1$  Column Vector

$$X_1 = F_{\text{act}}(W^T X_0)$$


$F_{\text{act}}$  and  $G_{\text{act}}$  =  
activation funcs that  
are applied to each  
tensor element. They  
do not alter its shape.

$Z$  = a  $2 \times 5$  Weight Matrix

$5 \times 1$  Column Vector



$$X_2 = G_{\text{act}}(Z^T X_1)$$

$$X_2 = G_{\text{act}}(Z^T F_{\text{act}}(W^T X_0))$$



# Bias Nodes

A bias node has a constant output of 1 but has K independent output weights, one per node in the layer. These weights are modified by gradient descent just like all other weights. However, they are often initialized to zero.



# Adding a Column Vector for Biases

- $B_w$  = a column vector of biases, one per node in the layer that produces  $X_1$ .
- $B_z$  = a column vector of biases, one per node in the layer that produces  $X_2$

$$X_1 = F_{act}(W^T X_0 + B_w)$$

$$X_2 = G_{act}(Z^T X_1 + B_z)$$

... or ....

$$X_2 = G_{act}(Z^T F_{act}(W^T X_0 + B_w) + B_z)$$

# Using a MiniBatch of Inputs

- Instead of an  $m \times 1$  column vector ( $X_0$ ) as input, use a minibatch ( $C$ ) of  $n$  column vectors.  $C = m \times n$  matrix.
- Every activation vector ( $X$ , originally a  $q \times 1$  column vector) now becomes a  $q \times n$  matrix ( $\chi$ ), one column per minibatch item.
- The  $q \times 1$  bias vectors ( $B$ ) become  $q \times n$  matrices ( $\beta$ ), where  $\beta = n$  copies of  $B$ . This is called **broadcasting**: expanding a matrix along one or more dimensions via copying.
- **The weight matrices and activation functions remain the same.**

$$\chi_1 = F_{act}(W^T C + \beta_w)$$


$$\chi_2 = G_{act}(Z^T \chi_1 + \beta_z)$$

... or ....

$$\chi_2 = G_{act}(Z^T F_{act}(W^T C + \beta_w) + \beta_z)$$






# Forward Propagation of a Minibatch

$$\begin{matrix}
 W^T & C \text{ (Size 4 Minibatch)} \\
 \left[ \begin{array}{|c|c|c|} \hline w_{11} & w_{21} & w_{31} \\ \hline w_{12} & w_{22} & w_{32} \\ \hline \end{array} \right] & \left[ \begin{array}{|c|c|c|c|} \hline \text{red} & \text{blue} & \text{orange} & \text{purple} \\ \hline \text{red} & \text{blue} & \text{orange} & \text{purple} \\ \hline \text{red} & \text{blue} & \text{orange} & \text{purple} \\ \hline \end{array} \right] & = W^T C \\
 2 \times 3 \text{ Matrix} & 3 \times 4 \text{ Matrix} & 2 \times 4 \text{ Matrix}
 \end{matrix}$$



Most tensor calculations, including those deriving gradients, work seamlessly with the extra tensor dimension due to the minibatch

# Classical Activation Functions






# The Sigmoid Activation Function

- Sigmoid (a.k.a. **logistic**) function **was** very popular for NNs, because a) common in biological systems, and b) implements a step/ramp but is continuous at the threshold, thus simplifying gradient calculations.
- But the sigmoid **saturates** for inputs of large (pos or neg) magnitude. These flat regions of the sigmoid curve have near-zero derivatives.
- Once a sigmoid saturates, the weights leading into that node have near-zero gradients  $\rightarrow$  they change very little  $\rightarrow$  learning halts.
- Hence, sigmoids are a poor choice for **hidden** nodes, particularly in deep networks.
- However, they are still very useful for **output** nodes, especially when combined with an objective function that reduces the risk of saturation (as discussed later).
- Also very useful in NNs that do not do backpropagation, and thus do not use gradients.
- The hyperbolic tangent ( $\tanh$ ) resembles the sigmoid, but since  $\tanh(0) = 0$  and it can also output negative values, it behaves a little more like an identity function (for small inputs). This makes it a bit easier to train.


# Contemporary Activation Functions

These have **greatly** improved the performance of NNs.


ReLU




Softplus




Leaky ReLU



Exponential Linear Unit (ELU)



# The Rectified Linear Unit (ReLU)



- Other activation functions yield error gradients that are strongly attenuated during backprop, such that only the last (closest to output) few layers of weights adapt properly.
- ReLUs permit sustained influence of gradients upon all weights throughout the network.
- By making many activation levels a hard zero, ReLUs **sparsify** activation patterns, which is more biologically realistic and better for learning.

Deep sparse rectifier neural networks, Glorot et. al., 2011

# Comparing Gradients: Sigmoid -vs- ReLU

- $v = V_j$  = sum of inputs to the activation function
- $x = x_j$  = output of the activation function

$f_T$  = Rectified Linear Unit (ReLU):  $f_T(v) = \max(0, v)$

$$\frac{\partial f_T(v)}{\partial v} = 1 \quad \text{when } v > 0$$

$$\frac{\partial f_T(v)}{\partial v} = 0 \quad \text{otherwise}$$

The gradient is **significant** for all  $v > 0$

$f_T$  = Sigmoid:  $f_T(v) = \frac{1}{1+e^{-v}}$

$$\frac{\partial f_T(v)}{\partial v} = x(1 - x)$$

- This has max = 0.25 (when  $x = 0.5$  and  $v = 0$ ). When the sigmoid saturates, it is much lower:  $\frac{\partial f_T(v)}{\partial v} = 0.048$  when  $x = 0.95$  (or  $x = 0.05$ ).
- These derivs get multiplied repeatedly (via Chain Rule) as backprop moves upstream, so products can get very small, very quickly.

## ReLU Problems and Solutions

- **Dying Neurons:** When sum of inputs ( $V_i$ ) is non-positive, the neuron outputs a zero AND its gradient is zero, so nothing changes. Inactive neurons often fail to come back to life.
- **Solution: Leaky ReLU:** Outputs may be small negative values, but these keep the neuron semi-active and yield non-zero gradients, which can eventually revive the neuron. Variations include Randomized Leaky ReLU (RReLU) and Parametric Leaky ReLU (PReLU).
- **Instability:** Gradients bounce around from positive to zero (when  $V_i \approx 0$ ) due to discontinuity of ReLU and Leaky ReLU at 0, where deriv from right = 1 but deriv from left = 0.
- **Solution: Exponential Linear Unit (ELU):** This is negative for  $V_i < 0$  (like the Leaky ReLU) but is continuous at  $V_i = 0$ , thus reducing oscillations of the gradient.

# Softmax - For Classification Problems

- Many classification problems benefit from an output layer that represents a probability distribution over the possible classes.
- Hence, all outputs must be non-negative, and they must sum to 1.
- No individual activation function can enforce this summation condition.
- Outputs must be combined and scaled (Boltzmann scaling).
- For an output layer of size  $N$  with a vector of output values,  $x_i$ :

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_{k=1}^N e^{x_k}}$$

- Note how this easily combines both positive and negative  $x_i$ , insuring that every softmax'd value is positive.
- Softmax embodies **competition** among the outputs, thus modelling lateral inhibition found in many brain regions.

# Scaling Outputs with Softmax

*Accentuating strength and weakness of output activations*

1.0	1.0	2.0	1.0	1.0
-----	-----	-----	-----	-----

Softmax



0.15	0.15	0.40	0.15	0.15
------	------	------	------	------

1.0	1.0	2.0	-1.0	-1.0
-----	-----	-----	------	------

Softmax



0.2	0.2	0.54	0.03	0.03
-----	-----	------	------	------

1	2	3	4	5
---	---	---	---	---

Softmax



0.01	0.03	0.08	0.23	0.64
------	------	------	------	------

-1	0	1	0	-1
----	---	---	---	----

Softmax









0.07	0.18	0.50	0.18	0.07
------	------	------	------	------

# Cross Entropy as a Loss Function

- Useful when targets and outputs both represent probability distributions. So combine softmax'd outputs with cross-entropy loss function.
- When combined with a sigmoid, the log counteracts the exp to reduce saturation (and gradient decay).
- The standard usage: targets = 1-hot vectors and outputs are a normalized distribution over class probabilities.

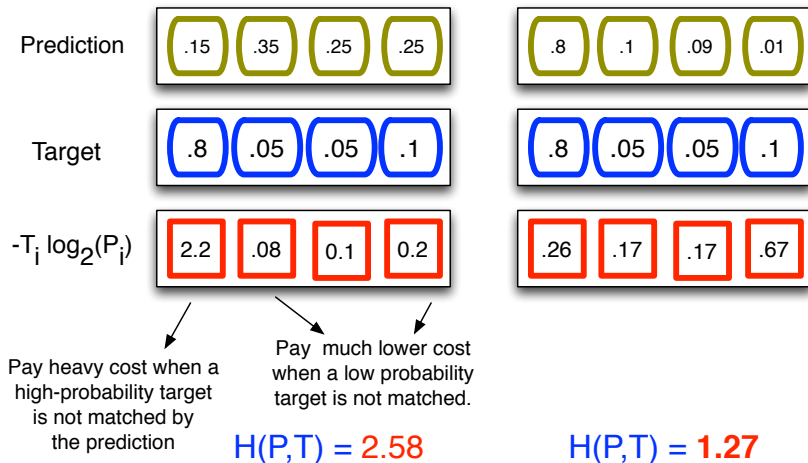
$$\text{Cross Entropy: } H(P, T) = - \sum_{k=1}^{\|T\|} T_k \log(P_k)$$

where P = network's output (prediction) and T = target vector.

Prediction		
One-hot target		
$-T_i \log_2(P_i)$		
Reduced cost when predictions match the hot bit.	$H(P, T) = 2.0$	$H(P, T) = 0.15$



# Cross Entropy to Compare 2 Distributions



# Cross Entropy for Logistic Regression


**Logistic (or Logit) Regression:** Single-value prediction where output = prob. of class membership. E.g. Insurance risk or not?

- The output node typically uses a sigmoid activation (a.k.a. logistic function), whose output value (in range (0,1)) represents the probability of class membership,  $p_k$ , for the  $k$ th case.
- Each target value,  $t_k$ , is binary.
- Cost function, called **Log Loss** applied to minibatch  $C$ :

$$L(C) = \frac{-1}{\|C\|} \sum_{k \in C} t_k \log(p_k) + (1 - t_k) \log(1 - p_k)$$


- Depending upon  $t_k$ , only one of the 2 terms is non-zero for each case,  $k$ :
  - $t_k = 1$  :  $p_k$  should be **high** to reduce cost.
  - $t_k = 0$  :  $p_k$  should be **low**, so  $(1 - p_k)$  is high, to reduce cost.
- Same principle as cross entropy, but now the single output node represents **two** classes: YES or NO. When the target = NO, the second term of the sum kicks in, asking, "How close is the prediction to a NO?"

# Cross Entropy Loss Overview



# Maximum Likelihood

Data Source



## The Modeling Objective

- Given a data source,  $S$ , build a model that best captures the statistical relationships (between features and classes) of  $S$ .
- View  $S$  as one big state, then the model ( $\theta$ ) should maximize the probability of that big state, i.e., the probability of ALL cases in the state being true:

$$\pi(S) = \prod_{k \in S} p(f_k, t_k) = \prod_{k \in S} p(t_k | f_k) p(f_k)$$

- where  $f_k$  = features and  $t_k$  = class (target) for case  $k$ .

# Conditional Maximum Likelihood

- The model ( $\theta$ ) will estimate  $p(t_k|f_k)$  with  $p_\theta(t_k|f_k)$ . E.g.  $\theta$  are the weights and biases of a neural network that takes  $f_k$  as input and produces an output vector of probabilities ( $P_k$ ), one for each class.
- $\theta$  and  $p(f_k)$  are independent, as the latter are determined solely by  $S$ .
- We want to choose  $\theta$  to maximize:

$$\pi_\theta(S) = \prod_{k \in S} p_\theta(t_k|f_k) p(f_k)$$

- This is equivalent to maximizing its log:

$$\log(\pi_\theta(S)) = \sum_{k \in S} \log(p_\theta(t_k|f_k)) + \log(p(f_k)) = \sum_{k \in S} \log(p_\theta(t_k|f_k)) + \sum_{k \in S} \log(p(f_k))$$

- The latter (red) sum is independent of  $\theta$ , so maximizing  $\pi_\theta(S)$  is equivalent to maximizing the former (blue) sum.
- The **Conditional Maximum Likelihood Estimator** is defined as:

$$\theta_{ML} = \operatorname{argmax}_\theta \sum_{k \in S} \log(p_\theta(t_k|f_k))$$

- I.e., the  $\theta$  that maximizes the log output values of nodes corresponding to the **correct** class for each case.

# Cond Max Likelihood and Cross Entropy

- Assume each case  $k \in S$  – or, more likely, a sample ( $\tilde{S}$ ) of  $S$  – is run through the model (with parameters  $\theta$ ) to produce  $P_k$ . Take the log of each output, yielding  $\log(P_k)$ , which is then compared to a 1-hot target vector ( $T_k$ ) that encodes the correct class. This comparison of outputs to targets is done via a vector dot product, which selects 1 value.
- Thus, for each case,  $T_k \bullet \log(P_k) = \log(p_\theta(t_k|f_k))$ , so:

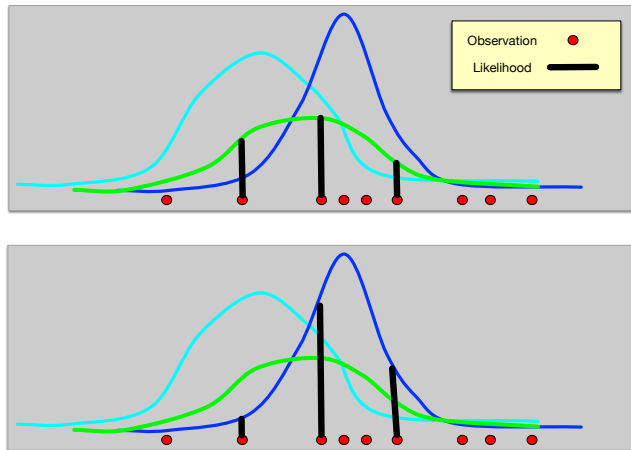
$$\sum_{k \in \tilde{S}} T_k \bullet \log(P_k) = \sum_{k \in \tilde{S}} \log(p_\theta(t_k|f_k))$$

- The same  $\theta_{ML}$  maximizes both, and minimizes the negation:

$$- \sum_{k \in \tilde{S}} T_k \bullet \log(P_k) = \text{cross-entropy}(T, P)$$

- Thus, the conditional maximum likelihood estimate ( $\theta_{ML}$ ) minimizes the cross entropy, a common objective function for classification problems. It also minimizes the mean squared error (MSE) for regression problems (details to follow).
- Cross entropy and MSE are very popular loss functions for DL.

# Max Likelihood of Observations



Goal: Find the parameters ( $\mu$ ,  $\sigma$ ) of a distribution (i.e. Gaussian / Normal) that maximizes the likelihood of all data/observations.

# Calculating Maximum Likelihood Estimate (MLE)

## The Normal (Gaussian) Distribution

$$p(x_i) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-(x_i - \mu)^2}{2\sigma^2}}$$

## Maximum Likelihood of a Gaussian

Given data  $(x_1, \dots, x_n)$ , MLE = values of  $\mu$  and  $\sigma$  that maximize this:

$$\prod_{i=1}^n p(x_i) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-(x_i - \mu)^2}{2\sigma^2}}$$

It's simpler, and equivalent, to find  $\mu$  and  $\sigma$  that maximize its log:

$$\begin{aligned} \log\left(\prod_{i=1}^n p(x_i)\right) &= \sum_{i=1}^n \log(p(x_i)) = \sum_{i=1}^n -\log(\sigma) - \frac{1}{2} \log(2\pi) - \frac{(x_i - \mu)^2}{2\sigma^2} \\ &= -n\log(\sigma) - \frac{n}{2} \log(2\pi) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 = \text{Log}L \end{aligned}$$



# Calculating MLE (2)

## Finding the MLE: $\hat{\mu}$ and $\hat{\sigma}$

Setting  $\frac{\partial \text{Log}L}{\partial \mu} = 0$  and solving for  $\mu$ , yields:

$$\hat{\mu} = \tilde{\mu} = \frac{1}{n} \sum_{i=1}^n x_i = \text{The sample mean}$$

Similarly

Setting  $\frac{\partial \text{Log}L}{\partial \sigma} = 0$  and solving for  $\sigma$ , yields:

$$\hat{\sigma} = \tilde{\sigma} = \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2} = \text{The sample standard deviation}$$

So the MLE for mean and standard deviation are just the mean and standard deviation of our original data !!

## Substituting $(\tilde{\mu}, \tilde{\sigma})$ for $(\mu, \sigma)$ in LogL

$$-n \log(\tilde{\sigma}) - \frac{n}{2} \log(2\pi) - \frac{1}{2\tilde{\sigma}^2} \sum_{i=1}^n (x_i - \tilde{\mu})^2 = -n \log(\tilde{\sigma}) - \frac{n}{2} \log(2\pi) - \frac{n\tilde{\sigma}^2}{2\tilde{\sigma}^2}$$

$$= -n \log(\tilde{\sigma}) - \frac{n}{2} (\log(2\pi) + 1)$$

# Calculating MLE (3)

To maximize  $\text{LogL}$ :

$$\text{LogL} = -n \log(\tilde{\sigma}) - \frac{n}{2} (\log(2\pi) + 1)$$

you need to minimize  $\tilde{\sigma}$ , where:


$$\tilde{\sigma} = \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \tilde{\mu})^2}$$

That's the same as minimizing the mean-squared error:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (x_i - \tilde{\mu})^2$$

⇒ Maximizing the Likelihood of Data = Minimizing that data's Mean Squared Error

# MLE and MSE for Regression



- $f$  = the regression function (blue)
- By maximizing the likelihood of the observations (by choosing a good function), we minimize  $MSE \propto \text{distance from observation } (x_i, y_i) \text{ (red point) to } (x_i, f(x_i))$ .
- For regression problems, MLE = the function that minimizes the MSE.